CSC411 Machine Learning
Project 3: Fake News

Ariel Kelman

Student No: 1000561368

Gideon Blinick

Student No: 999763000

19 March 2018

1 Dataset Description

Both datasets (real and fake) contain headlines about U.S. President Donald Trump. From the headlines, it is very much possible to determine with significant accuracy the percentage chance that a given headline is fake or real based on the precense of absence of certain words in the headline. For example, after performing some preliminary analyses on the headlines, we discovered that the following 3 words might be of particular use.

1. The word "donald" appears 42.05% of the time in real headlines while appearing in only 17.55% of fake headlines. This 24.5% difference was by far the largest of any word.

2. The word "the" appears 27.9% of the time in fake headlines while appearing in only 7.9% of real headlines for a difference of 20.02%.

3. The word "trumps" appears 11.1% of the time in real headlines while appearing in only 0.3% of fake headlines for a difference of 10.8%.

1.1 Results & Report Reproducibility

All results and plots can be generated using the code in `fake.py`. All code is python 3, run with Anaconda 3.6.3. Running the code will save all generated images in the `resources` folder, where they are used by LaTeX. Note that some of the sections require the code for other sections (in the same file) to be run first. To reproduce the report, simply run it through LaTeX. This will pull the most recently generated figures from the `resources` folder.

2 Naive Bayes Implementation

In part 2, we implemented the Naive Bayes algorithm on our training, validation and testing sets. To acheive this, we divided our solution into 3 functions and proceeded as follows:

1. In our first function "training_part" we use the training set to achieve estimates for p(real), p(fake), $p(word\|real)$ and $p(word\|fake)$. We estimated $p(word\|real)$ and $p(word\|fake)$ by counting the number of headlines in the training set where the word appears and the headline is

real, added a factor of $m \cdot p$, and divided this result by the size of the training set plus m. We also included words that were not found in the training set (but found in validation or testing) and set their values to $\dfrac{m \cdot p}{count(headlines) + m}$. The implementation of our function is"

```python
def training_part(fake_lines_training_set, real_lines_training_set, m, p):
    '''Takes in a list of training set lines divided into real and fake groups.
    Returns a list of probabilities that represent p(word|real) and p(word|fake).
    Therefore, this list will be 5833 long to account for every word.
    '''
    #Part 1: Training: calculate p(real), p(fake), and p(word|real), p(word|fake)
    fake_stats_training_set = get_stats(fake_lines_training_set) #compute probabilities for each word
    real_stats_training_set = get_stats(real_lines_training_set)

    fake_counts_training_set = get_count(fake_lines_training_set) #compute counts for each word
    real_counts_training_set = get_count(real_lines_training_set)

    p_fake = len(fake_lines_training_set)/(len(fake_lines_training_set) + len(real_lines_training_set))
    p_real = len(real_lines_training_set)/(len(fake_lines_training_set) + len(real_lines_training_set))

        # what we want is p(real|words in headline).
        # this is equal to p(words in headline|real)*p(real), divided by
        #   p(words in headline|real)*p(real) + p(words in headline|fake)*p(fake)
        # We already have p(real) and p(fake) from above, so we just need to find
        # p(words in headline|real) and p(words in headline|fake)
        # p(word in headline|real) is equal to count(word & real) + mp
        # divded by count(real) + m.
        # So we need to choose values for m and p and then come up with a
        # dictionaries of the adjusted probabilities.

    all_words = list(real_counts_total.keys())

    missing = { x:0 for x in fake_counts_total.keys() if x not in real_counts_training_set.keys() }
    real_counts_training_set.update( missing )
    missing = { x:0 for x in real_counts_total.keys() if x not in fake_counts_training_set.keys() }
    fake_counts_training_set.update( missing )


    adjusted_fake_counts_training_set = {}
    adjusted_real_counts_training_set = {}



    for word in all_words:
        adjusted_fake_counts_training_set[word] = fake_counts_training_set[word] + mp
        adjusted_real_counts_training_set[word] = real_counts_training_set[word] + mp

    naive_divisor = len(all_words) + m

    adjusted_fake_stats_training_set = {} #P(w | fake)
    adjusted_real_stats_training_set = {} #P(w | real)
```

```
    for word in all_words:
        adjusted_fake_stats_training_set[word] =
adjusted_fake_counts_training_set[word]/naive_divisor
        adjusted_real_stats_training_set[word] =
adjusted_real_counts_training_set[word]/naive_divisor


    return p_fake, p_real, adjusted_fake_stats_training_set,
adjusted_real_stats_training_set
```

2. In the second part , we created a function called "evaluate" which uses the outputs of the first function - p(real), p(fake), $p(word\|real)$, $p(word\|fake)$ - as its inputs, uses them to compute $p(headline\|real)$ and $p(healine\|fake)$ for a given set of headlines, and in turn uses those to compute $p(real\|headline)$ and $p(fake\|headline)$ for the set of headlines, which is the function output. In in this function that we utlize the fact that

$$a_1 \cdot a_2... = e^{loga_1 + loga_2...} \tag{1}$$

, in order to compute $p(headline\|real)$ as the product of $p(word\|real)$ for all words in and outside the statement. Performing the calculation in this way avoids underflow problems. The implementation of the function is:

```
def evaluate(p_fake, p_real, adjusted_fake_stats_training_set,
adjusted_real_stats_training_set, SET):
    ''' Calculate p(fake|headline) for all headlines in a given set (TRAINING,
TEST, VALIDATION).
    Takes p(real), p(fake), p(word|real), p(word|fake) as parameters.'''

    total_real_probabilities = {}
    total_fake_probabilities = {}

    all_words = list(real_counts_total.keys())

    for headline in SET:

        probabilities_real = []
        probabilities_fake = []

        headline_words = list(set( headline.split(' ') )) #converting to set
and back to a list removes duplicates
        for word in headline_words:
            probabilities_real.append(adjusted_real_stats_training_set.get(
word))
            probabilities_fake.append(adjusted_fake_stats_training_set.get(
word))
        non_headline_words = [x for x in all_words if x not in headline_words]
        for word in non_headline_words:
            probabilities_real.append(1 - adjusted_real_stats_training_set.get
(word))
            probabilities_fake.append(1 - adjusted_fake_stats_training_set.get
(word))

        total_real_probability = 0
        total_fake_probability = 0
        for k in range(len(probabilities_real)):
            total_real_probability = total_real_probability + math.log(
probabilities_real[k])
            total_fake_probability = total_fake_probability + math.log(
probabilities_fake[k])
        total_real_probability = math.exp(total_real_probability)
```

```
        total_real_probabilities[headline] = total_real_probability
        total_fake_probability = math.exp(total_fake_probability)
        total_fake_probabilities[headline] = total_fake_probability

    ## At this point, we have calculated p(headline|real) and p(headline|fake)
    for all
    ## headlines in whatever set we are testing
    ## Now we need to find p(fake|headline) which means multiplying p(headline
    |fake) by p(fake) and
    ## dividing by p(headline|fake)*p(fake) + p(headline|real)*p(real)

    #Compute final probabilites
    numerator_fake = [p_fake*x for x in total_fake_probabilities.values()]
    numerator_real = [p_real*x for x in total_real_probabilities.values()]

    final_fake = [numerator_fake[i]/(numerator_fake[i] + numerator_real[i])
    for i in range(len(numerator_fake))]
    final_real = [numerator_real[i]/(numerator_fake[i] + numerator_real[i])
    for i in range(len(numerator_real))]

    return final_fake, final_real
```

3. Finally, in the third part, we have a function called "check_accuracy" to determine the percentage of statements correctly classified. It does this by accepting $p(fake\|headline)$ from the previous function, as well as an array indicating whether a headline is fake by a 1 or real by a 0. The function then creates a new array that accepts $p(fake\|headline)$ for a given headline and rounds it to 1 if it is 0.5 or greater, and down to 0 otherwise. The function then subtracts this new array from our checker array and counts the number of non-zero elements. These elements represent incorrect classifications and allow us to determine our accuracy. Our function implementation is:

```
def check_accuracy(final_fake, y):
    '''Given a list of fake probabilities, compare with the actual results by
    rounding.
    If the fake  probability is greater than 0.5, consider it fake and if less
    , consider it real.
    Output the accuracy rate'''
    ## At this point we are done and have found p(fake|headline).
    ## All that remains is accuracy checking
    ## This part takes final_fake and one of y_tr, y_va, y_te as parameters
    ## and returns a percentage value in accuracy

    pred_fake = np.array([round(x) for x in final_fake])

    y_2 = np.array(y)

    incorrect = np.count_nonzero(pred_fake - y_2)
    total = len(y)
    correct = total - incorrect
    accuracy = correct/total

    return accuracy
```

In choosing our parameters m and p for the model, we began with a value of m $= 2 \cdot 5833$ and $\frac{1}{2 \cdot 5833}$ and optimized from there. This was chosen as our starting point because as mentioned on Piazza, it allows for 1 example per class per word and is intuitive. We then chose 4 different values of m and p and tested how they perform on our validation set, via a function called "optimize_mp". Our optimal m and p values were found to be 1 and 5833, respectively, achieving an acuracy on the validation set of 83.7%.

m	accuracy
5833	83.7%
$2 \cdot 5833$	83.1%
$3 \cdot 5833$	82.9%
$4 \cdot 5833$	82.7%

Our optimize_mp function is:

```python
def optimize_mp(fake_lines_training_set, real_lines_training_set, m_s, mp):
    val_acc = {}
    for m in m_s:
        print(m)
        p_fake, p_real, adjusted_fake_stats_training_set,
    adjusted_real_stats_training_set = training_part(fake_lines_training_set,
    real_lines_training_set, m, mp)
        final_fake, final_real = evaluate(p_fake, p_real,
    adjusted_fake_stats_training_set, adjusted_real_stats_training_set, validation_set
    )
        val_acc[m] = check_accuracy(final_fake, y_va)
    return val_acc
```

Finally, on the training set we obtained an accuracy of 83.4% while on the testing set we achieved an accuracy of 84.5%.

3 Predictive Factors

4 Logistic Regression

Code from project 2 was modified to run simple logistic regression to predict whether a headline is real or fake. For each headline, a sparse vector was created, with 1's in each position that represents a word that's present in the given headline (the representation of this vector is discussed again in section 5). The outputs were passed through the softmax function to produce results that can be interpreted probabalistically.

A regularization term was added directly to the gradient function that penalizes the L1 norm of the θ's (the parameters of the regression). To determine the best values for the learning rate and the weighting of the L1 regularization term (parameterized by γ), a simple grid search was preformed (primarily focussed on varying γ), using the results on the validation set to determine the best hyperparameters. The learning curves for many of these tests can be found in the `resources/part4` directory.

The following learning curves shows the results for the hyperparameter settings $\gamma = 1$ with a learning rate of 10^{-3}, though the results did not vary much with changes to γ.

These parameters gave a final training accuracy of XXX, validation accuracy of XXX, and testing accuracy of XXX.

5 Logistic Regression vs. Naive Bayes

When forming a prediction on a headline, both Naive Bayes and Logistic Regression effectively compute

$$\theta^T I = \theta_0 + \theta_1 I_1 + \theta_2 I_2 + ... + \theta_k I_k > thr$$

where thr is a given threshold value; in this project $thr = 0.5$. In both cases, the θ's represent the paremeters that define the prediction model, while the I's are a function of the input for which a prediction is produced.

Figure 1: Learning curve for the logistic regression with a learning rate of 10^{-3} and $\gamma = 1$.

For logistic regression, each $I = I(x)$ for a headline is simply a vector of 0's and 1's, where each position in the vector represents a word that was included in the training set. A 1 indicates that the word is found in the headline, while a 0 indicates that one was not. As the headlines are relatively short in comparison to the total number of words, this vector is very sparse. The θ's represent how important a feature (the presence of a particular word) is to the final prediction (θ_0 is a bias that can account for the priors on real vs. fake headlines).

For naive bayes,

6 Analysis of Logistic Regression

word	θ_i
hello	1

Table 1: table showing

7 Decision Tree

7.1 Classification

Using the `sklearn` implementation of decision trees, we trained several decision trees to differentiate between the fake and real headlines. After some experimentation with the many parameters in the `sklearn DecisionTreeClassifier` (particularly with the maximum number of features used when looking for the best split), the default parameters gave the best results. As a split condition, maximum entropy was used rather than gini impurity, though similar results were obtained for both. Many of the settings served as ways to limit the size of the decision tree, so this result is not unexpected.

Decision trees with a maximum depth of $\{2, 3, 5, 10, 15, 20, 35, 50, 75, 100, None\}$ were built, and the results on the training, validation, and testing sets are shown in the figure below. The final point, not plotted, with no limit on the maximum depth, gave an accuracy of $1, 0.760.76$ on the training, validation, and testing sets respectively. As can be seen from the figure, the larger the depth of the


Figure 2: Plot showing the preformance of `sklearn` decision trees with varying depth. The split condition was maximum entropy, there was no limit on the maximum number of features for a split, and stop words were included.

decision tree, the greater accuracy achieved on the testing set. However, improvement is small on the validation and testing sets after a depth of 20. Without any other constraints (such as a minimum number of samples to split a node), a decision tree can reach perfect accuracy on the training set, as demonstrated above. Predictably, this leads to much greater preformance on the training set than on the validation and testing sets (though there is no *decrease* in preformance on the latter sets); showing the importance of using a validation and testing set to measure preformance.

7.2 Visualization

The following image shows the first few layers of the decision tree with depth 20. It was generated by saving a text representation of the tree (as a `.dot` file), which was then visualized by using the

Figure 3: The top few layers of the decision tree (depth = 20).

It is interesting to look at the words being split on in these layers; X was a list of all words that appeared in either the fake or real datasets. These words are 'donald' (23), 'trumps' (1604; note the double appearance), 'the' (81), 'hillary' (44), and 'trump' (3). The exact same process of training was run ignoring stop words, but results were slightly worse; graphs similar to those mentioned above can be found in the `resources` directory (each filename indicates whether stop words were included, as well as info on the other non-default parameters). Note that because of the use of dictionaries (i.e. because `dict.keys()` does not keep track of order), running the code again may result in different indices for each of the above words.

A visualization of the entire tree is saves as `part7b_all.pdf` in the `resources` directory; the text represenations of many trees that were generated during training are saved in `resources/part7`.

7.3 Comparison of All 3 Classifiers

All three classifiers preform much better than random guessing...

8 Information Theory

8.1 Mutual Information on First Split

Using the result from part 7, the word that produced the best initial split was 'donald'. The mutual information (a measure of 'information gain') on that split (on the training data) can be calculated by

$$I(Y,x) = H(x) - H(x,Y) = H(Y) - H(Y,x)$$

where I is the mutual information, $H(x)$ is the entropy of x, and $H(x,Y)$ is the entropy of x conditional on Y. Plugging in the values shown in Figure 3 gives

$$\begin{aligned} I(Y,x) &= H(Y) - H(Y,x) \\ &= 0.969 - \left[P_1(1) + P_2(0.722) \right] \\ &= 0.969 - \left[\frac{1557}{2287}(1) + \frac{730}{2287}(0.722) \right] \\ &= 0.0557 \end{aligned}$$

where $x =$ '$donald'$ and Y indicates whether a headline is real or fake.

A function `mutual_info()` was written to check this result, and to calculate the mutual information for other possible splits. The code for this function is shown below; it gave a result of 0.58 for $x =$ '$donald'$.

```
def mutual_info(word, y):
    count_word = 0 #number of headlines with word in training set
    lines_with_word = []
    lines_without_word = []
    for k in range(len(training_set)):
        if word in training_set[k]:
            count_word += 1
            lines_with_word += [k]
        else:
            lines_without_word += [k]
    prob_word = count_word/len(training_set)

    y_with = np.array( [y[i] for i in lines_with_word] )
    y_without = np.array( [y[i] for i in lines_without_word] )

    prob_fake = np.count_nonzero(y)/len(y)
    H = prob_fake*np.log(prob_fake) + (1 - prob_fake)*np.log(1 - prob_fake) #
entropy before split
    H = - H/np.log(2) #convert to base 2, and apply negative

    prob_fake_with = np.count_nonzero(y_with)/len(y_with)
    Hy = prob_fake_with*np.log(prob_fake_with) + (1 - prob_fake_with)*np.log
(1 - prob_fake_with)
    Hy = - Hy/np.log(2) #entropy of headlines with word

    prob_fake_without = np.count_nonzero(y_without)/len(y_without)
    Hn = prob_fake_without*np.log(prob_fake_without) + (1 - prob_fake_without
)*np.log(1 - prob_fake_without)
    Hn = - Hn/np.log(2)   #entropy of headlines without word
```

```
        I = H − (Hy*len(y_with) + Hn*len(y_without) )/len(y)
        return I
```

8.2   Mutual Information on Later Split

The same procedure can be followed to compute the mutual information for a split on another word. Choosing a word randomly, gave a value of 0.0013 for the mutual information of a split of the training set on the word 'ways'. As expected, this value is lower than the mutual information for a split on the word 'donald' - that word was chosen as the first split precisely because it gave the most information about whether a headline was real or fake.